

INVESTOR PITCH DECK

AI Agent Security & Execution Integrity

We make it visible what AI agents are doing, why they are doing it, and when they must be stopped.

Action Visibility for AI Agents



PROBLEM → SOLUTION → PRODUCT → TECHNOLOGY → MARKET → GO-TO-MARKET

Would you give a brand-new employee access to every system in your company on day one?



Email



Payments



Database



Code



Cloud / API




CRM

**And this employee takes responsibility
for none of the mistakes it makes.**

Would you want that employee in your company?

EVIDENCE

This Is Already Happening



AUTHORIZATION $\neq$ CORRECTNESS

The Real Problem

Being authorized doesn't mean the action was right.

WHAT TODAY'S SYSTEMS CHECK

- Who is this agent?
- Is it authorized to use this tool?
- Is the request technically valid?

WHAT THEY DON'T CHECK

- ? Why is it doing this?
- ? What triggered this decision?
- ? Is it consistent with previous steps?
- ? Does the action have irreversible impact?

89%

Can't get agents into
production


11% make it through · Source: KPMG AI
Pulse 2026

The agent's identity is verified. Its chain of reasoning is not.

VERIFYING THE DECISION CHAIN

Our Solution

“Baleena evaluates an agent’s action history, its authority and the trustworthiness of the data it used together. Every critical action is classified with a risk score: blocked, routed for human approval, or verified after execution.”



TECHNICAL DEPTH

Product & Technology

Protection & Decisioning

- Real-time risk engine
- Policy-based decisions, blocking and approval
- State- and history-aware behavior auditing

Traceability & Trust

- Why-Chain cause-and-effect analysis
- Authority and data-provenance tracking
- Privacy-preserving telemetry

Action Security

- Verifies that the action an agent wants to take matches what the user actually asked for.
- Blocks risky actions and detects unexpected outcomes.

AI Agent 1 / AI Agent 2

Sends request



Baleena Broker

Policy Check

Risk Scoring

Why-Chain




Approve → Real endpoint

Block → Held

SCENARIO-BASED SELECTION · ARCHITECTURAL APPROACH

Decision Model

The agent chooses from a catalog of pre-approved scenarios. A decision model makes that choice as a proposal, not a decision.



01 · UNTRUSTED BY DESIGN

The model proposes; only the policy engine decides.

02 · NARROWING ONLY

Scenarios can tighten permissions — never widen them.

OUR DIFFERENTIATOR: BOTH AT ONCE

Token Efficiency & Predictability

FIXED RULE SETS

Predictable, but every situation must be written as a rule in advance.

SCENARIO CATALOG + DECISION MODEL

Fewer tokens and higher predictability, together.

The agent selects from a closed set instead of reasoning from scratch at every step: less token use, a shorter decision path, and outcomes bounded by pre-approved scenarios.

FREE AGENT REASONING

Flexible, but reasons from scratch at every step. More tokens, longer decision paths, less predictable outcomes.

Token limits or token governance alone don't deliver this: the combination does.

Architectural approach under evidence-gated validation; qualitative comparison.

 **Baleena**
Execution integrity

Command

Overview

Actions

Investigate

Execution graph

Trust and data flow

Sessions

Reviews

Discover

Inventory

Risk and coverage

Govern

Policy studio

EffectSpec studio

Remediation

Decision model

Operate

Deployments

System health

Audit

RUNTIME LINK

CONNECTED

bundle: bundle@2026-08-12.1
planner/v1.1.0
mode: OBSERVE

Command

Overview

Observe

Enforce

HYBRID

acme - staging

updated just now

Decision against execution

Two separate verdicts. Neither one implies the other.

Decision activity

Sep 28, 2026, 22:04:28 – Sep 28, 2026, 22:10:32

500 decisions the planner produced

ALLOW

REVIEW

DENY

324

30

146

Requirements nobody could establish

0 UNKNOWN

Execution

Inspect allowed actions

294

Actually dispatched to a tool

Policy allowed 324. A decision is not an execution – the two counts differ, and the difference is below.

0 Allowed but never run

Every allowed action in this window reached dispatch.

NOT EXECUTED not a DENY – policy did not refuse it

29 Dispatched, still unresolved

The exposure is held and reconciliation has not resolved it.

IN DOUBT not a failure – the exposure stays held

Live telemetry

Filtered execution stream.

View in ledger →

Q Filter by tool or action...

Any decision

Any attention

8 recent actions

	Decision time	Agent / Subject	Action	Policy decision	Disposition	Settlement	Reason
	Sep 28, 2026, 22:10:29	travel-agent	refund_execute REFUND	ALLOW	NOT EXECUTED	not dispatched	–
	Sep 28, 2026, 22:10:29	travel-agent	refund_execute REFUND	ALLOW	EXECUTE	CONFIRMED	–
!	Sep 28, 2026, 22:10:29	travel-agent	refund_execute REFUND	DENY	NOT EXECUTED	not dispatched	authority.ceiling_exceeded
!	Sep 28, 2026, 22:10:29	travel-agent	refund_execute REFUND	REVIEW	PENDING REVIEW	not dispatched	approval.required
	Sep 28, 2026, 22:10:29	travel-agent	refund_execute REFUND	ALLOW	EXECUTE	CONFIRMED	–
	Sep 28, 2026, 22:10:29	travel-agent	refund_execute REFUND	ALLOW	EXECUTE	CONFIRMED	–
	Sep 28, 2026, 22:10:27	travel-agent	refund_execute REFUND	ALLOW	EXECUTE	CONFIRMED	–
	Sep 28, 2026, 22:10:27	travel-agent	refund_execute REFUND	ALLOW	EXECUTE	CONFIRMED	–

A bounded in-process window over recent decisions. It does not survive a restart and is not the audit record. Actions beyond the window were observed and then dropped from this view, not from the journal.

Product · Actions Feed, live action monitoring

10 / 19

Baleena

Execution integrity

Command

Overview

Actions

Investigate

Execution graph

Trust and data flow

Sessions

Reviews

Discover

Inventory

Risk and coverage

Govern

Policy studio

EffectSpec studio

Remediation

Decision model

Operate

Deployments

System health

Audit

RUNTIME LINK

CONNECTED

bundle bundle@2026-08-12.1

plannex/v1.1.0

mode OBSERVE

Command > Actions

← All actions

Decision record

req-795-9d14a2878718 · customer_data_export on customer:{customer_id} for tenant acme, decided Sep 28, 2026, 22:11:04.

POLICY DECISION

• DENY

what the planner produced

→

COMMIT RESULT

not reached

whether the scarce right could be taken

→

EXECUTION DISPOSITION

• NOT_EXECUTED

what the runtime did

→

SETTLEMENT

not dispatched

what it later learned about the effect

Four separate facts in the order they occur. None is derived from another: an ALLOW whose guard was refused is NOT_EXECUTED and never a DENY, and a dispatch whose outcome never came back is IN_DOUBT rather than failed.

policy denied the action

Primary reason: INTERPRETATION_AMBIGUOUS

DUPLICATE_KEY/customer_id

Linked to requirement: interpretation.integrity

Decision spine

Each stage the runtime walked, in order, with what it established and how far the action got. Not a why-chain: that is the derivation over recorded occurrences, on the execution graph.

Request

OK

Interpret

DEGRADED

Identity

OK

Authority

OK

Effect

DEGRADED

State, session and sequence

UNKNOWN

Invocation

BLOCKED

Influence

BLOCKED

Provenance

DEGRADED

Enforcement

DEGRADED

Degradation

DEGRADED

Interpret

DEGRADED

the body was not strictly parsed, not complete, or carried an ambiguity

PARSING ASSURANCE

STRICT

COMPLETENESS

COMPLETE

BINDING

CANONICAL_REWRITE

DECLARED LEVEL

REDUCED

SCHEMA FINGERPRINT

sha256:8e77019_538735358a

BODY MEDIA TYPE

application/json

BODY BYTES

149

AMBIGUITIES

DUPLICATE_KEY

DETAIL

/customer_id

field is supplied by both the canonical route and request body

Product · Action Detail, risk scoring and intervention

11 / 19



Baleena

Execution integrity

Command

Overview

Actions

Investigate

Execution graph

Trust and data flow

Sessions

Reviews

Discover

Inventory

Risk and coverage

Govern

Policy studio

EffectSpec studio

Remediation

Decision model

Operate

Deployments

System health

Audit

Runtime link

CONNECTED

bundle: bundle@2026-08-12.1

planner/v1.1.0

mode: OBSERVE

Historical

Proposed

ALLOW

REVIEW

DENY

Observe

Enforce

HYBRID

acme · staging · v

updated 10s ago

Policy studio

The bundle in force, a proposed draft, and what the draft would have decided about decisions that already happened. A proposal is not a decision, and a counterfactual is not an execution.

Current bundle

bundle@2026-08-12.1 · 2 configurable rules · evaluated by planner/v1.1.0

refund_execute.authority_ceiling

DENY

authority_ceiling authority.ceiling

ceiling_known && amount_minor > ceiling_minor

refund_execute.high_value_review

REVIEW

threshold_review approval.high_value

amount_known && amount_minor > 20000

No control plane signs this bundle yet. The canonical, closure-free artifact a signature would cover is assembled and held, which is not the same as a signature existing.

Proposed draft

Local to this browser. Compiled by the runtime's own loader when you run the counterfactual; never saved, never activated.

BUNDLE SOURCE (YAML)

own message.

#

Nothing here is saved. This console has no policy write API and does not

create one; the draft lives in this browser until you leave the page.

version: "proposed@draft-1"

rules:

- id: refund_execute.high_value_review

tool_id: refund_execute

kind: threshold_review

condition: "amount_known && amount_minor > 5000"

decision: REVIEW

requirement: approval.high_value

reason_matched: approval.required

always_record: false

Run counterfactual

Reset draft

Activation is not offered. There is no policy write API on this surface.

Policy is authored in the bundle and compiled through the validator and the CEL type checker before it can be evaluated. This console reads it; it does not author it.

Counterfactual against recorded history

bundle@2026-08-12.1 → proposed@draft-1 · 200 of 200 retained decisions replayed

All retained

Market Analysis

TOTAL ADDRESSABLE MARKET

\$10-15B

Global AI cybersecurity / AI governance / agent runtime security market — growing 45% a year to \$90–100B by 2030

SERVICEABLE ADDRESSABLE MARKET

\$250-400M

→ \$1B MENA, Türkiye and Europe; companies running agents in production

SERVICEABLE OBTAINABLE MARKET

\$50M

Obtainable market size at 5% capture of SAM

Focus: AI agents taking real actions in high-volume customer operations across MENA, Europe and Türkiye

B2B SaaS · BPO/CCaaS · Digital Travel · E-commerce · Fintech

INTERVIEWS

30+

Customer interviews completed

LETTERS OF INTENT

5

Letters of intent signed

STATUS

Ready

Customers ready for PoC

Sources: Precedence Research; project customer & market analysis; KPMG AI Pulse agent deployment trend. With no official data set, organisation counts are a segment-level market assumption: 600–1,800 organisations × \$20–50K annual contract value.

SELF-HOSTED · PAID POC + ANNUAL / MONTHLY LICENSE

Revenue Model

01 · PAID POC

Pilot / proof phase

A paid proof of concept on the customer's own environment — the entry point to a license.

02 · ORGANIZATION LICENSE — PRIMARY & GROWING REVENUE

Annual or monthly license

One license per organization. Web self-service purchases carry the same annual license price — not a separate tier.

SELF-HOSTED DEPLOYMENT

The customer hosts the deployment and carries all infrastructure cost. No pricing by usage volume, agent count or action count.


INCLUDED IN THE LICENSE

Token-usage governance — tracking and limits per agent, user and department — and the decision model. No add-on tier.

Early-stage pricing assumption; license terms will follow customer scale and deployment.

OUR DIFFERENCE: NOT PROMPT SECURITY — ACTION-CENTRIC RUNTIME GOVERNANCE

Competitive Landscape



OUR CATEGORY

Execution Assurance

- **Authority Assurance:** Authority/trust propagation, delegation chain
- **Behavior Assurance:** Stateful obligation tracking, small classifiers
- **Impact Assurance:** Agent action transaction: stage/commit/verify

The only position here offering token savings and higher predictability at the same time — via scenario-based selection.

HOW TO READ THIS

- Independent agent-security vendor
- Large security platform
- Cloud-native platform substitute
- Observability / adjacent (indirect)

Bubble size = commercial maturity · Early → Scaled / funded

● Baleena (orange) — token savings + predictability together

Team


AMAZON STARTUP ACTIVE



EV LAB


**Nilsu Demirağ**Founder — Research &
Design**Birdem Üstündağ**Founder — Product &
Engineering**Yavuz Selim Yaşar**Founder — Business
Development**Arif Emre Kılıç**Founder — Backend &
Security**Cansın İsmail Bahçeci**Founder — Product &
Operations**ACADEMIC ADVISOR****Assoc. Prof. Dr. Taner Arsan**

AI, Cybersecurity, Network Engineering, Anomaly Detection

**TECHNOLOGY & INDUSTRY MENTOR****Uğur Kazdal**AWS Solutions Architect MENAT — Cloud Architecture,
Microservices, Kubernetes · 10+ years of experience

FOUNDER-LED SALES + SELF-SERVICE DEMO + TECHNICAL CONTENT

Roadmap



Financial Projection

Month 18

Break-even Target

12

Active customers in year 3

\$600K

Year 3 revenue target

\$140K

Year 3 net profit target

23%+

Year 3 operating margin

ASSUMPTIONS

PoC/demo: free, limited feature set

Year 1: 4 early customers

Year 2: 8–10 active customers

5-year financial projection (target, \$ thousands)





**We make your AI agents auditable
and safe action by action without
slowing them down.**

Thank You

info@getbaleena.com